







(a) Restoration optimisation: classical OR baselines, RL surrogate, and learned-surrogate MPC are statistically equivalent in this regime (N=5 seeds for stochastic methods)

(b) Planner-relevant ranking metric: MSE-only $\approx$ contrastive in restoration

(c) Ranking failure is domain-specific: present in farmland, absent in restoration

(d) Buchanan Pareto front (7 weight configs): risk $\leftrightarrow$ water trade-off

(e) Synthetic Pareto front: shape parallels Buchanan

(f) Buchanan trajectory: contrastive MPC vs greedy heuristic (shapes match $\rightarrow$ both reach near-optimum on this case)

(g) Reward decomposition: contrastive vs MSE near-identical

