

Reproducible learned-surrogate planning for county-scale farmland-consolidation scenario audit

Ning Zhou^{a,b,**} and Xiang Jing^b

^aPublic R&D Center, SuperMap Software Co., Ltd., Beijing 100015, China

^bSchool of Software and Microelectronics, Peking University, Beijing 100871, China

Abstract

Farmland consolidation is a spatial-planning lever for mechanisation, slope management, and large-patch formation in mountainous counties, but county-scale scenario search is difficult because cadastral action spaces contain tens of thousands of parcels and rewards are multi-objective. We present a learned transition-model and model-predictive-control pipeline for rapid candidate generation with independent GIS audit. On Bishan District, the unconstrained shipped scenario is trained in 25 min and planned in ~ 5 min on a desktop CPU; GIS recomputation reports a -2.001% farmland-slope change and a 687 ha shift out of the steep ($\geq 15^\circ$) bands, but also a 505 ha cultivated-area loss. Adding a hard no-net-loss cultivated-area floor converts the audit finding into an execution constraint: Bishan still reduces slope by -0.675% , raises contiguity by 0.020, gains 29.8 ha of qualifying *baimu fang*, and increases total cultivated area by 3.0 ha; a seven-profile Bishan execution frontier exposes the slope–area trade-off, with no-net-loss profiles retaining -0.608 to -0.711% slope change. Neijiang Dongxing similarly moves from a -261 ha cultivated-area loss to $+62.2$ ha. On the matched research-side cadastre used for in-house centralised-PPO and multi-agent reinforcement-learning baselines, contrastive MPC reaches $-1.289 \pm 0.079\%$ slope change across five independently trained ensembles, $1.6\times$ either DRL baseline; against simulated annealing under a matched 180 s wall budget the reward gap is $9.6\times$. A pairwise ranking loss explains when learned-surrogate planning helps, and a public-data restoration case bounds its transfer. The workflow ships as an ArcGIS Pro toolbox and CLI, with open reproduction tracks.

1 Introduction

Mountainous and hilly landscapes contribute a substantial share of global agricultural production but pose management challenges distinct from contiguous large-field plains: small, irregular cadastral parcels distributed across steep terrain raise per-hectare labour cost, prevent the use of medium-scale machinery, increase vulnerability to climate-driven yield shocks, and limit policy levers for raising productivity without displacing farmers [Romeo et al., 2020, Liu and Li, 2023, Niu et al., 2023, Foley et al., 2011, Springmann et al., 2018]. Land consolidation—the spatial rearrangement of farmland and forest parcels within legal and topographic constraints—is the dominant policy response [Liu et al., 2014, Long, 2014, Rounsevell et al., 2012]. In China, where a substantial fraction of cultivated land lies in hilly or mountainous counties [Huang et al., 2018], consolidation is operationalised through the formation of *baimu fang* (literally “hundred-*mu* patches”, ≥ 6.67 ha of contiguous farmland) and through township-level regulatory plans that bound feasible swaps between farmland and forest classes.

^{**}Corresponding author. zhouning1@supermap.com

Computationally, county-scale consolidation is a high-dimensional, multi-objective combinatorial planning problem: a single county routinely contains tens of thousands of parcels, three competing objectives (slope reduction for mechanisation, contiguity for shape, large-patch consolidation for management), and hard adjacency and tenure constraints. Classical approaches—rule-based heuristics, genetic algorithms [Cao et al., 2012, Deb et al., 2002, Stewart et al., 2004, Aerts and Heuvelink, 2002], and model-free deep reinforcement learning [Zheng et al., 2023, Equihua et al., 2024, Janowicz et al., 2020]—all face structural limits at this scale: heuristics ignore the multi-objective trade-off, metaheuristics miss long-horizon dependencies, and model-free RL in our experience requires 8–12 GPU-hours per run and fails on 25–40% of random seeds in the discrete county-scale regime, costs that prevent iterative scenario exploration by planning offices.

Model-based reinforcement learning offers a structurally appealing alternative: learn a transition model once from data, then exploit it for cheap planning [Chua et al., 2018, Janner et al., 2019]. Recent work in continuous control [Hafner et al., 2025, Hansen et al., 2024] and game-playing [Schrittwieser et al., 2020, Silver et al., 2017] demonstrates the paradigm; AI-for-sustainability and Earth-observation work has begun to adopt these methods at policy scale [Rolnick et al., 2022, Reichstein et al., 2019, Tuia et al., 2024]. We use *learned planning model* for a feed-forward deterministic next-state predictor with a discriminative reward head, rather than the recurrent latent “world model” of the Dreamer family, because our setting is fully observed and episodic on pre-extracted cadastral block features.

Contribution. This manuscript reports three contributions, in decreasing order of generality. First, we present an *empirical characterisation* of learned-surrogate model-predictive control as a rapid scenario generator for county-scale spatial planning, anchored in two real Chinese counties (Bishan, Chongqing, 52,515 parcels; Neijiang Dongxing, Sichuan, 76,376 parcels) plus a public-data US restoration case (Buchanan County, Virginia, 562 units). We deliberately separate optimisation output from policy compliance: the Bishan package-side candidate reduces area-weighted farmland slope by -2.001% under independent GIS-only recomputation, shifts 687 ha out of the $\geq 15^\circ$ slope bands, and raises contiguity by $+0.013$, but the same audit shows a 505 ha cultivated-area loss. We then impose the audited red-line requirement as a hard cumulative cultivated-area floor and show that no-net-loss constrained runs remain productive on both real counties (Bishan: -0.675% slope, $+0.020$ contiguity, $+29.8$ ha *baimu fang*, $+3.0$ ha cultivated area; Neijiang: -0.492% slope, $+0.037$ contiguity, $+252.6$ ha *baimu fang*, $+62.2$ ha cultivated area). Second, we package the method comparison as two empirical diagnostics that a practitioner can evaluate before adopting the pipeline: the within-state-to-between-state reward variance ratio σ_a/σ_s predicts whether an auxiliary pairwise ranking loss is required to make standard mean-squared-error (MSE) training planning-useful, and simulator step cost combined with reward-interaction structure predicts whether a learned surrogate is worth its training overhead. We frame both diagnostics as calibrated on the present evidence, not as universal thresholds. Third, we release a reproducibility-oriented toolchain: trained ensembles, a deterministic seven-preset synthetic benchmark under CC-BY 4.0, the public-data Buchanan restoration case, an ArcGIS Pro Python Toolbox, a non-commercial CLI workflow, and a ~ 70 -second smoke test. On the research-side cadastre that matches the pipeline on which our in-house centralised-PPO and multi-agent RL baselines were trained, five independently trained contrastive ensembles reach $-1.289 \pm 0.079\%$ slope change, $1.6\times$ the magnitude reached by either DRL baseline (§4, Welch’s t -test, $|d| > 1.9$), while reducing wall time by roughly two orders of magnitude. The auxiliary pairwise large-margin loss [Burgess et al., 2005, Joachims, 2002] that operationalises the first diagnostic raises within-state ranking accuracy from 51.6% to 85.5% and converts model-predictive control from the worst variant tested to the best. Pairwise-preference

training itself is not new [Joachims, 2002, Burges et al., 2005, Christiano et al., 2017, Ouyang et al., 2022]; the contribution is its empirical role, auditability, and boundary in spatial planning.

2 Multi-objective consolidation as a planning problem at county scale

We anchor the analysis in two complete county-scale datasets from southwest China: Bishan District in Chongqing (13 townships, 52,515 parcels grouped into 2,600 spatial blocks) and Neijiang Dongxing District in Sichuan (29 townships, 76,376 parcels grouped into 3,711 blocks; per-county initial-state metrics are reported in Supplementary Table S1). Together the two counties comprise 128,891 individual cadastral parcels, an order of magnitude larger than published consolidation benchmarks we are aware of [Cao et al., 2012, Zheng et al., 2023]. Area-weighted mean farmland slope is 9.62° on Bishan and 10.55° on Neijiang—both above the 6° terracing threshold of the Chinese High-Standard Farmland Construction standard (GB/T 30600-2022) beyond which mechanised cultivation requires terraced fields. The within-farmland contiguity index is 3.59 on Bishan and 2.63 on Neijiang. The legally relevant consolidation unit, the *baimu fang* (≥ 6.67 ha contiguous farmland eligible for mechanised management), is unevenly distributed: 109 patches totalling 46,844 ha on Bishan, 384 patches totalling 74,342 ha on Neijiang.

Three properties make the problem hard. The action space is combinatorial—across a 100-step episode the joint decision space is approximately 2600^{100} on Bishan alone, vastly larger than any standard combinatorial benchmark in the spatial-planning literature. The objectives interact: slope-reducing swaps near fragment boundaries often retire marginal farmland from the largest existing patches, so a planner who optimises slope greedily can lose qualifying area, and vice versa. The wall-time profile of existing methods rules out iterative scenario exploration: in-house Centralised PPO and multi-agent baselines on the same task require 8–12 GPU-hours of training per seed and Centralised PPO collapses on 2/5 seeds (Supplementary Section S1). A policy-compliant plan must reduce farmland slope, raise contiguity, increase or preserve useful large-patch structure, and satisfy cultivated-area controls under paired farmland–forest swap and township-budget constraints. In the experiments below, the learned planner generates and audits candidate scenarios; it does not by itself prove that every unconstrained candidate is acceptable for implementation.

3 Diagnosing and fixing a hidden ranking failure in the standard learned model

We trained a learned planning model—an ensemble of three transition networks ($\sim 237,000$ parameters each, mean-squared-error objective)—on 6,000 transitions sampled from random and greedy policies in the Bishan environment. By conventional fidelity metrics the model is excellent: state cosine similarity exceeds 0.998, training and validation losses converge cleanly, and the model is rolled forward autoregressively for 25 steps before predictions diverge perceptibly (s -cosine drops from 0.999 at one step to 0.94 at 50 steps).

The aggregate metrics, however, conceal a near-total failure of per-action discrimination. We constructed a pairwise-action benchmark by snapshotting the simulator state at each of 1,000 sampled environment states, executing 50 valid actions one at a time, recording the immediate ground-truth reward, and restoring the snapshot before continuing. The MSE-trained ensemble correctly ranks only 51.6% of action pairs by predicted reward—statistically indistinguishable from random—while predicting a reward standard deviation across the action menu of 0.274 versus the

ground-truth 0.811, a three-fold compression of the discriminative signal. A fresh re-train with the identical objective reaches 60.2% ranking accuracy on the same benchmark; both points sit well below the 85.5% that the auxiliary ranking loss subsequently delivers (Table 1). We therefore treat 51.6% as the lower envelope of MSE-only behaviour at this σ_a/σ_s regime rather than its central tendency, and the 25–34 pp gap between MSE-only and contrastive ensembles—rather than the absolute level reached by any single MSE-only retrain—as the empirical fingerprint of the diagnostic. Four independent operational failures follow at the deployment layer: dream-trained policies collapse under uncertainty-guided evidence collection because 99.6% of candidate transitions are filtered out by an ensemble-disagreement signal that carries no actual signal (the standard ensemble disagreement diagnostic shows Pearson $r = -0.024$ between disagreement and absolute prediction error); the random-sampling MPC ceiling is $-0.209 \pm 0.008\%$; every aggressive-exploitation variant fails catastrophically (greedy continuation gives $+0.029\%$, the wrong sign, and slope-only scoring gives -0.001%); and doubling the candidate budget to $K=100$ provides no improvement.

The decomposition behind this is structural to the loss. In plain terms, two quantities govern whether MSE training delivers a planning-useful reward predictor. The between-state spread σ_s measures how much average reward varies from one state to the next; the within-state spread $\sigma_a(s)$ measures how much reward varies between the candidate actions available at a single state. Planning depends almost entirely on the latter—top- K selection compares actions inside a single state, never across states—but MSE training is dominated by the former whenever $\sigma_a(s)$ is small relative to σ_s , because the loss-minimising predictor in that regime is the conditional mean $\bar{r}(s)$ which assigns the same reward to every action regardless of identity. On Bishan, the empirical ratio is $\sigma_a/\sigma_s \approx 0.6$, the regime where the conditional-mean predictor gives random ranking by construction. Aggregate prediction fidelity (cosine 0.998) and within-state ranking accuracy (51.6%) are therefore not constrained against each other under MSE training, and the latter is what model-predictive top- K selection actually consumes; we give the formal statement and its evaluation-methodology implications in Methods (§9).

The diagnosis points at a concrete intervention: add an explicit pairwise large-margin ranking term [Herbrich et al., 2000, Joachims, 2002, Burges et al., 2005, Burges, 2010] to the model’s reward head, applied alongside the existing reconstruction objective (Methods, §9):

$$\mathcal{L} = \underbrace{\mathcal{L}_{\text{state}} + \mathcal{L}_{\text{global}} + 0.1 \mathcal{L}_{\text{reward}}}_{\text{MSE terms}} + \lambda_{\text{rank}} \mathcal{L}_{\text{rank}}. \quad (1)$$

The signal is applied to the existing reward head, not a separate reward proxy, so deployment requires no additional model. This is a model-side use of pairwise preferences, distinct from the policy-side RLHF formulation [Christiano et al., 2017, Ouyang et al., 2022]. Training data are inexpensive: 1,000 states sampled by random rollout, 50 valid actions snapshotted/executed/restored at each, generated in ~ 45 min on commodity CPU.

The intervention closes the discriminative gap with no measurable cost to aggregate fidelity (Table 1). Pairwise ranking accuracy rises from 60.2% ($\lambda_{\text{rank}}=0$) to 85.5% ($\lambda_{\text{rank}}=5.0$) while state-prediction cosine similarity stays above 0.999. The qualitative effect on MPC is a reversal: greedy continuation, the variant most cleanly attributable to ranking error, becomes the best variant, and cross-seed mean slope improvement on Bishan reaches $-1.289 \pm 0.079\%$ across five independently trained ensembles, an approximately six-fold improvement over the MSE-only ceiling of -0.209% .

Table 1: Auxiliary contrastive training restores per-action discrimination and converts MPC from worst variant to best. Top: ranking diagnostics on 1,000 held-out action pairs (predicted-reward standard deviation across the 50-action menu, pairwise rank accuracy, state-cosine fidelity); the true reward standard deviation across actions (0.811) bounds the recoverable signal. Bottom: MPC slope on real Bishan ($H=5$, $K=50$). Single-ensemble $\lambda \in \{0, 1.0\}$ rows and the $\lambda=5.0$ random row report episode-level standard deviation; the starred $\lambda=5.0$ greedy row reports cross-seed mean across 5 independently trained ensembles. A fresh MSE-only re-train ($\lambda=0$) only reaches 60.2% rank accuracy, confirming the failure is structural to the objective rather than to the optimiser. Absolute pairwise-accuracy values depend on the evaluation pairwise distribution; the qualitative monotonic trend (rank accuracy rises with λ_{rank}) reproduces on independently-generated pairwise datasets, and we release the trained checkpoints together with an evaluation script (§Code Availability).

<i>Ranking diagnostics on held-out pairwise data</i>				
Configuration	λ_{rank}	Pred. reward std	Rank accuracy	Cos sim
Baseline ensemble (cached, MSE only)	—	0.274	0.516	0.998
Retrained (MSE only)	0.0	0.715	0.602	0.9997
Contrastive (weak)	1.0	0.389	0.816	0.9997
Contrastive (strong)	5.0	0.183	0.855	0.9991
True reward std (reference)	—	0.811	—	—
<i>MPC outcome on Bishan ($H=5$, $K=50$)</i>				
λ_{rank}	Continuation	Slope %	vs. baseline	Status
0.0 (MSE only)	random	-0.291 ± 0.020	-0.082	marginal
0.0 (MSE only)	greedy	-0.252 ± 0.007	-0.042	marginal
1.0	random	-1.310 ± 0.033	-1.101	strong
1.0	greedy	-1.376 ± 0.010	-1.167	strong
5.0	random	-1.344 ± 0.049	-1.135	best random
5.0*	greedy	-1.289 ± 0.079	-1.080	best overall (5 seeds)
Baseline ceiling (MSE, random)	random	-0.209 ± 0.008	—	reference
Baseline (MSE only, greedy)	greedy	$+0.029 \pm 0.005$	$+0.238$	catastrophic

4 Real-county deployment: Bishan and Neijiang

The clearest test of the intervention is whether it can generate candidate planning scenarios that are useful enough to inspect, fast enough to iterate, and explicit enough to reject when they violate policy constraints. We deployed the contrastive-MPC pipeline end-to-end on Bishan District, Chongqing (52,515 parcels, 2,600 blocks at the research-side cadastre; 53,004 swappable parcels at the package-side prepare on which we report the independent audit) and recomputed all primary outcomes directly from the optimised cadastral shapefile written to disk by an independent GIS-only script that shares no source with the optimisation pipeline (Methods §9). Figure 1 therefore summarises an unconstrained candidate scenario rather than a final plan: panel (a) shows the per-step slope-reduction trajectory of the shipped contrastive ensemble against the in-house Centralised PPO and MARL baselines, and panel (b) shows the spatial distribution of the approximately 850 paired farmland–forest swaps committed on the optimised cadastre. Per-seed multi-objective profiles for the lab-trained five-ensemble cohort that anchors the DRL comparison are reported in Supplementary Table S9; the on-disk verification, model-audit, policy-audit, constrained no-net-loss numbers, and Bishan execution-constraint frontier are reported in Supplementary Tables S10–S17.




Figure 1: Bishan unconstrained candidate scenario and independent audit. (a) Per-step slope-reduction trajectory of the shipped Bishan ensemble on the package-side prepare, mean of five evaluation episodes; the ensemble is deterministic across the action-selection path under fixed seed, so the five episodes overlap to within 0.002 pp at every step. Endpoint annotations report the env-side slope (-2.04%) and the independent GIS-only recomputation of -2.001% (Methods §9). The dashed line shows the in-house Centralised PPO baseline ($-0.79 \pm 0.36\%$, lab pipeline) and the dotted line shows MARL ($-0.81 \pm 0.10\%$, lab pipeline); the open marker at the corresponding endpoint is the lab-pipeline contrastive-MPC result of $-1.289 \pm 0.079\%$ (5 ensembles $\times$ 1 episode each, Supplementary Table S9). The DRL baselines were trained on the lab pipeline; the package-side value is not collapsed with the lab-pipeline value because the corrected geographic-CRS slope step exposes truer initial slopes (9.62° vs 8.41°). (b) Spatial distribution of the paired swaps committed by the planner on the optimised cadastre (episode 0), shown as farm→forest parcels (red, slope-shed) and forest→farm parcels (green, slope-take), buffered to 250 m for visibility at county scale. Swaps concentrate at fragment boundaries where slope-reducing trades cost the least qualifying area. This candidate is useful for audit, but it is not yet a final policy-compliant plan because it reduces total cultivated area by 505 ha (Supplementary Table S14).

The independent audit shows why the scenario is informative but not automatically deployable.

On the package-side prepare, area-weighted mean farmland slope decreases by 2.001% relative to the initial 9.62°, within-farmland contiguity rises by 0.0128, and the count of qualifying *baimu fang* (≥ 6.67 ha contiguous farmland units) increases by +8 patches relative to the initial 109. The slope change is concentrated in policy-relevant bands: optimised farmland gains +48 ha in the $< 6^\circ$ band and +134 ha in the 6–15° band, at the expense of –615 ha from the 15–25° band and –72 ha from the $> 25^\circ$ band, a total of 687 ha shifted out of the steep tail. However, the same audit identifies two costs. Total qualifying *baimu fang* area decreases by –473 ha, equivalent to 1.01% of the initial 46,844 ha qualifying stock. More importantly for policy compliance, the paired farm $\leftrightarrow$ forest swap conserves parcel *counts* but not parcel *areas*: the planner sheds an aggregate 843 ha of high-slope farmland and admits 337 ha of lower-slope forest, yielding a net cultivated-area decrease of 505 ha (1.05% of the 48,283 ha initial farmland stock). We therefore treat the Bishan package-side output as a constraint-audited candidate scenario. It shows that the learned surrogate can expose a high-value slope-contiguity trade-off quickly, but it also shows that cultivated-area preservation must be encoded as a hard execution constraint before administrative use. Across the 13 Bishan core townships (Supplementary Table S15), per-township slope improvement ranges from –0.57% to –4.08% ($\mu = -2.07\%$, $\sigma = 1.09\%$, Gini = 0.28), indicating that the candidate is not concentrated in a single township even though it is not red-line-compliant.

We therefore re-ran the deployment pipeline with a hard cumulative cultivated-area floor of 0 ha, so every committed farm $\rightarrow$ forest / forest $\rightarrow$ farm pair must keep total cultivated area at or above the initial stock (Supplementary Table S16). This no-net-loss constraint reduces Bishan’s slope-improvement frontier but preserves policy feasibility: the independently recomputed output reaches –0.675% slope change, +0.0204 contiguity, +1 qualifying *baimu fang*, +29.8 ha qualifying *baimu fang* area, and +3.03 ha total cultivated area after 401 paired swaps. The constrained run still shifts 218 ha out of the $\geq 15^\circ$ steep tail, so the slope benefit is not a metric artefact of shrinking cultivated area.

To characterise the policy trade-off rather than report a single repair point, we also ran a seven-profile Bishan execution-constraint frontier (Supplementary Table S17 and Supplementary Figure S1). Allowing a cultivated-area loss of roughly 500 ha retains nearly the unconstrained slope gain (–1.746% versus –1.753% on the default released ensemble) but remains outside a no-net-loss rule. Tightening the floor to –250 and –100 ha yields slope changes of –1.261% and –0.918% while reducing steep-tail farmland by 469 and 329 ha. The three no-net-loss profiles, including a conservative profile that also forbids loss of qualifying *baimu fang* area and a connectivity-conservative profile, all satisfy the cultivated-area floor and retain non-trivial slope improvements (–0.608 to –0.711%), positive contiguity changes (+0.0204 to +0.0259), and positive qualifying-area changes (+29.2 to +43.4 ha). This frontier changes the actual committed swaps through hard execution constraints and local pair-selection settings; it is not presented as a full retrained reward-weight Pareto sweep.

Comparison against model-free reinforcement-learning baselines requires a matched cadastre, because the in-house Centralised PPO and multi-agent variants (Supplementary Section S1) were trained on the lab pipeline (8.41° initial farmland slope) before the prepare-step correction described in Methods §9. We therefore re-ran five independently trained contrastive ensembles end-to-end on the lab pipeline and report this as the baseline-comparison regime, not as the package-side deployment audit. Slope change under contrastive MPC ($-1.289 \pm 0.079\%$, 5 ensembles $\times$ 1 episode each on the lab cadastre) is approximately $1.6\times$ larger in magnitude than the multi-agent baseline ($-0.81 \pm 0.10\%$; Welch’s $t = -8.01$, $df=7.9$, $p < 10^{-4}$, Cohen’s $d = -5.07$) and the Centralised PPO baseline ($-0.79 \pm 0.36\%$; Welch’s $t = -3.01$, $df=4.5$, $p = 0.034$, Cohen’s $d = -1.90$), with comparable contiguity (+0.016) and *baimu fang* count (+3.4 versus +3.2 and +2.9). Cross-seed slope variance is $4.6\times$ smaller than Centralised PPO and similar to MARL’s. The total qualifying-

area loss on the lab pipeline is -312 ± 34 ha (versus -75 to -97 ha for DRL baselines), confirming that the stronger slope objective comes with a larger area trade-off. The cost asymmetry is the main operational result of this comparison: the contrastive ensemble trains in 25 minutes on a commodity 12-thread CPU and runs each 100-step planning episode in approximately 5 minutes, against 8–12 GPU-hours per seed for the in-house DRL baselines.

We also evaluated classical operations-research baselines on Bishan at a 180s wall budget matched to the contrastive-MPC pipeline: simulated annealing over a permutation of block selections reached cumulative episode reward 9.04 at 100 iterations, random-restart over 98 episodes reached 30.85, and contrastive MPC reached 86.74—a $9.6\times$ gap over simulated annealing. Extending the OR baselines to a $10\times$ longer 1,800s wall budget (Supplementary Section S12) raises simulated annealing only to reward 9.54 at 1,187 iterations (+5.5%) and random-restart to 37.70 at 1,105 episodes (+22%), an order of magnitude below the contrastive-MPC comparison value under the tested budgets. The 180-s comparison is reported as a single matched-budget run per algorithm rather than as a seeded distribution because the SA and random-restart baselines are deterministic-by-construction at fixed RNG and therefore degenerate as a between-seeds comparison; for the contrastive-MPC pipeline, the cross-seed reward standard deviation in the open package-side reproduction (5 ensembles, Supplementary Section S3) is $\sigma(R) = 7.4$ at the comparable operating point, an order of magnitude below the reward gap. The mechanism (the simulator’s 17.9ms/step cost combined with combinatorial cost interactions in the reward function) is developed in §7.

The intervention reproduces on Neijiang Dongxing (Sichuan, 76,376 parcels, 3,711 blocks, contiguity 2.63 versus Bishan’s 3.59; five seeds re-run from scratch). Slope decreases by $0.50 \pm 0.02\%$, contiguity rises by 0.034 ± 0.001 , and the *baimu fang* area trade-off flips sign to a $+267 \pm 38$ ha gain. A matched GIS-only policy audit of the canonical unconstrained Neijiang output nevertheless finds the same cultivated-area failure mode as Bishan, though smaller in relative terms: slope decreases by -0.664% , contiguity rises by $+0.032$, but total cultivated area decreases by 261 ha (0.334%), because 880 ha of farm→forest parcels are offset by only 619 ha of forest→farm parcels. Under the same hard no-net-loss floor, Neijiang retains most of the policy-relevant gains while satisfying the area constraint: slope decreases by -0.492% , contiguity rises by $+0.0373$, *baimu fang* area increases by 252.6 ha, and total cultivated area increases by 62.2 ha after 466 paired swaps. Three differences from Bishan reflect terrain morphology rather than algorithmic instability: slope improvement is roughly half (Neijiang’s flatter mean leaves less slope-reduction headroom); contiguity improvement is roughly twice as large (the more fragmented start has more room to consolidate); and the constrained area trade-off is positive because consolidation operates on small fragments rather than retiring patches from the largest existing ones. Cross-seed slope standard deviation ($\pm 0.02\%$) is an order of magnitude smaller than the Bishan in-house DRL baselines.

Two cadastral pipelines, kept separate. The Bishan numbers cited in this section come from two internally consistent pipelines and should not be read as a single pooled effect size. The *package-side prepare* (53,004 swappable parcels, initial slope 9.62° under a corrected geographic-CRS slope step, Methods §9) is the public-facing deployment cadastre that ships with the toolchain, on which the independent GIS audit of the unconstrained candidate scenario is computed. The *lab pipeline* (52,515 parcels, initial slope 8.41° , no geographic-CRS correction) is the cadastre on which the in-house Centralised PPO and multi-agent baselines were trained, and on which we therefore reran the contrastive ensembles for the matched-cadastre comparison ($-1.289 \pm 0.079\%$, 5 ensembles $\times$ 1 episode each). The ~ 0.7 pp gap between the two slope-change endpoints is attributable to the corrected prepare step, which exposes truer initial slopes and therefore leaves more reduction headroom. A complete cross-pipeline reconciliation table, including the $-1.544 \pm 0.041\%$ five-

episode toolchain-CLI re-run and the per-seed open-source -1.738% end-to-end reproduction, is reported in Supplementary Section S3 so readers can audit the evidence chain explicitly.

Robustness and auditability. Four checks bound the Bishan and Neijiang results. (i) A direct GIS recomputation from the optimised cadastral shapefile written to disk reproduces the slope, contiguity, *baimu fang*, and cultivated-area deltas to floating-point noise without invoking the learned ensemble (§9, package-side pipeline); both the unconstrained violation and the constrained no-net-loss repairs are therefore properties of the on-disk geometry, not of the learned dynamics. (ii) The contrastive ensemble’s 1-step prediction error is measured directly against the true environment along the same 100-step trajectory MPC visits (§9); the ensemble is mis-calibrated as a regressor (reward MAE 0.82 against $\sigma(r_{\text{true}}) = 1.40$) yet recovers a Spearman rank correlation of $+0.22$ sufficient at $H=5$ to reproduce the action sequence to floating-point identity. (iii) A counter-factual ablation that replaces the ensemble with the true `env.step` inside the scoring loop returns a slope of -1.819% in 358 s, worse on all three reported objectives than the package-side ensemble pipeline’s -2.001% in 1,448 s under the implemented compute concessions (§9). (iv) Cross-county replication on Neijiang shows the qualitative slope-contiguity result survives a $1.4\times$ scale-up and a different morphology, and the same cultivated-area-floor mechanism removes the red-line failure in both counties.

5 Cross-landscape stress test on a synthetic benchmark

To test whether the result depends on properties shared by Bishan and Neijiang (both situated in southwestern hilly terrain), we constructed a synthetic benchmark of seven parametric landscapes varying along three morphological axes: terrain (hilly, mixed, plain), problem size (800, 2,600, 4,500 blocks), and fragmentation (consolidated, fragmented). Two presets are calibrated against the real anchors; the remaining five are parameter-controlled stress tests (Methods, §9). All seven landscapes plus the deterministic generator are released under CC-BY 4.0 (§Data Availability), so external researchers can reproduce our experiments and benchmark future methods without access to restricted cadastral records.

We benchmarked five methods (uniform-random selector, slope-greedy heuristic, genetic algorithm [Cao et al., 2012], centralised PPO, contrastive MPC) on the seven landscapes with five random seeds each (full slope-improvement profile in Supplementary Table S5). Two patterns mark the boundary of the claim. On the five presets that are large or fragmented enough to make multi-step planning informative (the two real-data clones plus three medium-large fragmented presets), contrastive MPC is the strongest method, with PPO the consistent runner-up; the MPC slope-improvement margin over the genetic algorithm is 6–54% across these cells, and remains 37% even on the largest preset (4,500 blocks, ~ 9.3 h MPC wall time vs ~ 248 s GA). On the small consolidated preset (`plain_small_cons`, 800 blocks), the genetic algorithm wins ($-2.757 \pm 0.203\%$ vs. MPC $-2.226 \pm 0.220\%$): the action space is small enough that GA’s population-based exploration covers it within budget, and the planning value of contrastive ranking, which scales with the within-state action menu, is muted. The synthetic stress tests therefore confirm the regime structure read from §4–§6 rather than overturn it: contrastive MPC is the strong method when the problem is large or fragmented enough to make per-state ranking matter; classical search wins when it is not.

6 Cross-domain verification on a public-data restoration case

To probe the boundary of the contrastive intervention beyond cadastral consolidation, we apply the unmodified pipeline to a structurally different geospatial planning problem: ecological-restoration prioritisation for abandoned mine lands. We test two cases. The first is a real public-data instance from Buchanan County, Virginia, USA: $n_{\text{units}} = 562$ planning cells at 2km resolution, derived from the OSMRE e-AMLIS public abandoned-mine inventory, USGS NHD flowlines, USGS 3DEP slope at 30 m, and Census TIGER county boundaries. The second is a deterministic synthetic instance with $n_{\text{units}} = 420$. Reward in both is a linear combination of risk reduction, water-quality protection, spatial connectivity, and cost penalty (Methods, §9); per-action reward is a per-unit attribute lookup.

We evaluate seven methods on each case: random selector, greedy by priority score, simulated annealing, NSGA-II, CBC-MILP weighted-sum, MSE-only learned MPC, and contrastive learned MPC (Table 2). Two findings emerge. First, on both cases the contrastive intervention provides no operationally meaningful advantage: on Buchanan, the contrastive ($\lambda=5$) and MSE-only ($\lambda=0$) ensembles agree to within 0.20 reward units (paired $t = 2.61$, $\text{df}=4$, $p = 0.060$; Cohen’s $d_z = 1.17$), and a two-one-sided-tests (TOST) equivalence test rejects a difference larger than ± 0.5 reward units ($\sim 0.22\%$ of mean reward 229.3) at $p_{\text{TOST}} = 0.008$, supporting equivalence within a pre-specified operational margin; pairwise ranking accuracy is ≥ 0.95 for both, and top- K regret is < 0.003 at $K=1$ for both. The decomposition of §3 explains this: in farmland the within-state action variance is small relative to between-state variance ($\sigma_a/\sigma_s \approx 0.6$ on Bishan), so MSE training collapses to the conditional-mean predictor; in restoration the per-action reward is a per-unit attribute lookup, $\sigma_a/\sigma_s \approx 5\text{--}7$, and MSE training acquires the ranking signal as a by-product. The contrastive intervention is therefore beneficial only when ranking failure is actually present. Second, classical OR baselines (greedy, simulated annealing, CBC-MILP) reach near-optimum in the restoration regime, with cumulative-reward gaps to the learned-MPC variants on the order of the cross-seed standard deviation of the learned variants ($\sigma_{\text{seed}} = 0.03$ for $\lambda=5$, 0.16 for $\lambda=0$, $n = 5$ each). The order-of-magnitude lift is over the random baseline, not over OR baselines: contrastive MPC more than doubles the random reward on Buchanan, but the gap to a well-engineered OR baseline is within seed-level noise. This is in contrast to farmland (§4), where simulated annealing reaches reward 9.04 versus contrastive MPC’s 86.74 at matched 180 s wall budget—a $9.6\times$ gap. The factor that toggles between these two regimes is the per-step simulator cost; we develop the full crossover analysis in §7.

7 Simulator step cost as the dominant predictor of where MPC beats classical OR

The cross-domain analysis of §6 closed one question (when does the contrastive ranking loss help?) but raised another. On the restoration cases, classical OR baselines reached near-optimum at the wall-clock budgets reported in Table 2, while on farmland the same algorithms underperformed by an order of magnitude. The natural answer is the cost of one simulator call: we test it directly by instrumenting the restoration env with a fixed sleep injected on every `env.step` and re-running simulated annealing and learned-MPC at matched 30-second wall budgets across delays of $\{0, 1, 10, 50, 100\}$ ms per step on each of ten (case $\times$ reward-profile) cells (Methods, §9).

The pattern across the cells is consistent with two independent factors. First, on nine of ten cells where reward is dominated by per-unit attribute lookups, the MPC-vs-SA gap is monotonic in step delay (Buchanan/delayed crosses zero at ~ 30 ms and reaches $+31.2\%$ at 100 ms): SA’s it-




Figure 2: Cross-domain verification, expanded baseline panel, and Pareto front. (a) Cumulative episode reward across seven methods on Buchanan (left axis) and synthetic (right axis); MSE-only and contrastive learned-MPC variants are statistically indistinguishable from greedy-by-priority and CBC-MILP, all far above the random baseline. (b) Top- K regret as a function of K : MSE-only and contrastive surrogates are equivalent on both cases, confirming that the contrastive intervention has no headroom to improve in the restoration regime. (c) Pairwise ranking accuracy across all four domains: farmland (Bishan, Neijiang) shows the large MSE-vs-contrastive gap, restoration (Buchanan, synthetic) shows MSE-only already saturating. (d, e) Pareto front of cumulative risk reduction vs cumulative water/habitat protection across seven reward-weight configurations on Buchanan and synthetic respectively; the shape is consistent across cases. (f) Trajectory comparison on Buchanan: per-step reward (bars) and cumulative reward (line) for contrastive MPC vs the greedy heuristic; both reach the same end state via similar action sequences, confirming that the OR baseline is a strong reference on this case. (g) Reward-component decomposition: contrastive vs MSE-only allocate cumulative reward across the four objective terms within noise.

Table 2: Optimisation results across seven methods on the two restoration cases (cumulative episode reward; mean $\pm$ std for stochastic methods, $n = 5$ seeds where applicable; deterministic methods are single runs). The greedy heuristic, simulated annealing, and CBC-MILP weighted-sum baselines reach near-optimum on both cases, statistically indistinguishable from the learned-surrogate MPC variants. The advantage of the learned-surrogate pipeline over the OR baselines is therefore not the optimum quality on this case; it is the wall-clock and the ability to absorb new reward-weight configurations without re-running the optimiser (§9).

Method	Buchanan VA (real, $n = 562$)			Synthetic (deterministic, $n = 420$)		
	total reward	n_{sel}	wall-time (s)	total reward	n_{sel}	wall-time (s)
random (uniform)	$+88.7 \pm 12.7$	50.0	< 0.1	$+6,163 \pm 160$	59.8	< 0.1
greedy by priority	$+228.76$	50	< 0.1	$+8,590.9$	60	< 0.1
simulated annealing	$+228.76$	50	0.7	$+8,737.0$	60	0.6
NSGA-II (4-obj)	$+175.45$	50	3.2	$+7,075.7$	60	2.9
CBC-MILP weighted-sum	$+229.43$	50	3.0	$+8,127.2$	49	< 0.1
MSE-only MPC ($\lambda=0$)	$+229.20 \pm 0.14$	50	~ 28	$+8,173 \pm 66$	60	~ 5
contrastive MPC ($\lambda=5$)	$+229.39 \pm 0.03$	50	~ 28	$+8,128 \pm 63$	60	~ 5

Algorithm-advantage crossover: simulator step cost determines whether learned-surrogate MPC beats classical OR (SA)




Figure 3: Simulator step cost is the dominant predictor of when learned-surrogate MPC beats simulated annealing under matched wall-clock budgets. (a) Reward gap (MPC minus SA, normalised by per-cell mean reward) versus injected step delay (log scale) on ten (case $\times$ reward-profile) cells. The shaded green band marks the empirical MPC-advantage zone observed on these ten cells (transition near 30 ms/step); we report this as an empirical inflection rather than a universal threshold. The actual farmland Bishan datapoint (gold star), measured at 17.9 ms/step natural step cost under a 180s wall budget, sits at +89.6%. (b) SA iteration count completed in the same wall budget on a log-y axis: the count collapses linearly with step cost from $\sim 9,000$ at $5 \mu\text{s}/\text{step}$ to ~ 6 at 100 ms/step. Below ~ 100 iterations SA cannot converge, the mechanism by which step cost translates into the MPC advantage. The synthetic/scale_economy line is the exception: MPC dominates at every step delay because its combinatorial cost-interaction structure is what learned-surrogate batched scoring exposes naturally.

eration count under a fixed wall budget falls from $\sim 9,000$ at $5\mu\text{s}/\text{step}$ to ~ 6 at $100\text{ms}/\text{step}$, while learned-MPC’s per-step cost ($\sim 25\text{ms}$ one batched ONNX forward pass over $K=50$ candidates) is decoupled from `env.step`. Below ~ 100 iterations the SA neighbourhood walk cannot converge on a 50-unit permutation problem. Second, the synthetic/scale_economy exception (MPC dominates SA by +542 to +931 across every step delay including $5\mu\text{s}$) shows that combinatorial cost interactions are an independent factor: when an action’s effective cost depends on the already-selected set, the joint interaction structure is not visible to a permutation-search procedure that ranks units independently, but the learned surrogate’s batched scoring at the current state internalises it naturally.

These two factors compose in farmland. The bishan CountyLevelEnv has a measured natural step cost of $17.9 \pm 0.3\text{ms}$ (5-step probe $\times 3$) and a reward function with combinatorial cost interactions (*baimu fang* count is union-find aggregation over per-block selections, contiguity is area-weighted adjacency average, slope is the area-weighted geometric integral that updates with every paired swap). Under matched 180s wall budgets, simulated annealing reaches reward 9.04 ($n_{\text{iter}} = 100$, near the convergence threshold), random restart reaches 30.85 ($n_{\text{eps}} = 98$), and contrastive MPC reaches 86.74. The empirical diagnostic the analysis suggests for problem selection is direct: report simulator step cost (in ms/step) and whether the reward function exhibits combinatorial cost interactions; in our measurements the learned-surrogate pipeline is in the advantage regime at the farmland operating point ($17.9\text{ms}/\text{step}$, combinatorial reward) and outside it at the restoration operating point ($0.05\text{ms}/\text{step}$, attribute-lookup reward), and the cross-cell injected-delay sweep in Figure 3 traces a transition near $30\text{ms}/\text{step}$ on the cells we tested. We treat 30ms as the empirical inflection observed on these ten cells rather than as a universal threshold; the mapping from step cost to MPC advantage is mediated by the SA iteration count that fits in the wall budget, which itself depends on the simulator’s per-step parallelisability and the reward landscape’s local-search difficulty. Whether the same inflection transfers to physics-driven Earth-science simulators (hydrology, power-grid maintenance, chemical-reaction prioritisation, supply-chain equilibria) is a hypothesis we discuss in §9 rather than a result we establish here.

8 From paper to planner’s desk: an end-to-end deployment toolchain

The intervention’s value to consolidation policy depends on whether a county planner—typically working in a desktop GIS environment with limited compute and no machine-learning support staff—can actually run it on local cadastral data. We packaged the full pipeline as a four-tool ArcGIS Pro Python Toolbox: **Prepare Data & Blocks** (Tool 1) ingests a standard Third National Land Survey DLTB shapefile and a digital elevation model and emits the block-level state expected by the learned-model code; **Sample Transitions** (Tool 2) generates the pairwise dataset $\mathcal{D}_{\text{pw}}$ on the user’s data; **Train Contrastive Ensemble** (Tool 3) trains the transition-model ensemble; and **MPC Planning** (Tool 4) runs the planning loop and writes back an optimised cadastral shapefile readable by any standard GIS workflow (Figure 4). A fifth diagnostic tool validates the user’s environment before the pipeline is allowed to run.

The runtime profile reflects a deliberate split between one-time setup (Tools 1–3, $\sim 1\text{--}2\text{h}$ on a 12-thread CPU per county) and repeated planning (Tool 4, $\sim 7\text{min}$ per scenario). An end-to-end UI-driven run on full Bishan, with the same training data and hyperparameters as Section 4, yields a slope change of $-1.544 \pm 0.041\%$ across five evaluation episodes of the shipped Bishan ensemble (same direction and order of magnitude as the lab figure of $-1.289 \pm 0.079\%$; the gap reflects two replication regimes, 5 ensembles $\times$ 1 episode in the lab versus 1 ensemble




Figure 4: The four-tool ArcGIS Pro deployment pipeline (§9). Tools 1–3 run once per county to prepare cadastral data, sample the pairwise dataset, and train the contrastive transition-model ensemble (exported to ONNX). Tool 4 is the iterated planning loop a county officer runs repeatedly under different reward weightings; one 100-step scenario finishes in ~ 7 min on a desktop CPU. A separate Tool 5 (dashed) validates the Python environment, Spatial Analyst availability, and ONNX shape compatibility before any pipeline run is allowed.

$\times 5$ episodes in the toolchain). To support users without an ArcGIS Pro license, the same end-to-end pipeline is also exposed through a non-commercial command-line interface (`farmland-mpc prepare/sample/train/plan`) that runs against any QGIS-compatible vector and raster stack via `geopandas/rasterio`. The package is released under MIT licence; the synthetic benchmark (Section 5) under CC-BY 4.0. Together with the trained Bishan ensemble and the released restoration cases (§6), this forms a reproduction kit for method behaviour and independent output auditing, while raw real-county cadastral records remain under the data-governance limits described in Data availability.

9 Discussion

What the pipeline delivers for sustainability planning. For consolidation policy in mountainous-terrain counties, the operational shift is from slow training runs to fast, auditable scenario generation. Model-free reinforcement-learning baselines previously required 8–12 GPU-hours per scenario and collapsed on roughly a quarter of attempted runs; classical operations-research methods at the same wall-clock budget reach less than 11% of contrastive MPC’s reward on Bishan. With contrastive MPC, a planner obtains a candidate scenario in minutes on a desktop, and the independent GIS audit can then expose whether that candidate satisfies slope, contiguity, large-patch, and cultivated-area constraints. The Bishan and Neijiang cases illustrate why this audit loop matters: unconstrained high-slope candidates are useful for understanding the slope-retirement frontier, but both lose cultivated area (505 ha and 261 ha respectively) and should therefore be treated as rejected scenarios. The hard cultivated-area-floor variant converts that audit finding into an executable constraint and preserves positive slope, contiguity, and large-patch gains in both counties. The Bishan execution-constraint frontier further shows how a planning office could move along explicit policy thresholds: each stricter cultivated-area floor gives up part of the maximum slope reduction, but the no-net-loss profiles remain beneficial on all audited objectives.

Methodological mechanism. Two mechanisms are stacked behind these gains, both characterised in §3 and §7 respectively. The auxiliary pairwise large-margin term closes a hidden ranking failure of the standard MSE objective in the $\sigma_a/\sigma_s \lesssim 1$ regime farmland inhabits, and is redundant when reward is a per-unit attribute lookup. The learned-surrogate pipeline as a whole, separately, beats classical operations-research baselines when each simulator step is computationally expensive (the empirical inflection in our injected-delay sweep is near 30 ms/step on the ten cells tested;

farmland operates at 17.9 ms/step naturally and at 1.8s per SA *iteration* once the combinatorial reward is taken into account) or when reward exhibits combinatorial cost interactions, or both. This mechanism explains the computational advantage; policy compliance still comes from explicit feasibility constraints and GIS audit, not from reward optimisation alone.

Practical implication for problem selection. Practitioners considering a learned-surrogate planner on a new domain should report two quantities: simulator step cost (in ms/step) and whether reward exhibits combinatorial cost interactions. In our measurements, the learned-surrogate pipeline is in the advantage regime at the farmland operating point (17.9 ms/step, combinatorial reward) and outside it at the restoration operating point (0.05 ms/step, attribute-lookup reward); on the injected-delay sweep across ten (case $\times$ reward-profile) cells the transition occurs near 30 ms/step, which we report as the empirical inflection on these cells rather than as a universal threshold. The diagnostic is calibrated on the present farmland and restoration evidence only. Applications with high-fidelity physics simulators, such as hydrologic-model-driven restoration or infrastructure maintenance scheduling, are plausible targets for future testing rather than established domains of validity.

Bounds on the claim. Five constraints qualify the result. First, the unconstrained Bishan and Neijiang package-side scenarios are not red-line compliant because total cultivated area decreases by 505 ha and 261 ha respectively; the constrained no-net-loss runs fix this specific failure, and the Bishan execution-constraint frontier maps seven feasible or infeasible policy profiles, but a full reward-weight Pareto study would still require re-sampling transitions and re-training the exported reward head under each weight setting. Second, cadastral evidence comes from two counties in southwestern China within the same regional regime; structurally different agricultural settings remain untested as far as the contrastive intervention is concerned. Third, the real-county raw cadastral records cannot be redistributed, so the fully open reproduction path uses the synthetic farmland benchmark and the public Buchanan restoration case, while the real-county path relies on released derived features, pairwise datasets, model checkpoints, and audit scripts. Fourth, a finer characterisation across the variance ratio σ_a^2/σ_s^2 would help practitioners decide a priori whether to add the auxiliary loss. Fifth, the pairwise dataset relies on simulator snapshot/restore, available in our gridded simulators but not in many physics-based environments. The immediate follow-up is therefore not another proof that cultivated-area preservation can be enforced, but a broader cross-county constrained-frontier and retrained reward-weight study.

Methods

Scope and architectural framing. The architecture used throughout this paper is a feed-forward deterministic learned transition model with a discriminative reward head (§9 below). It is in the operational sense of Chua et al. [2018] and Janner et al. [2019] a learned $\hat{p}(s_{t+1}, r_t \mid s_t, a_t)$ that substitutes for the real environment in MPC lookahead and dream-style policy training; it is not a generative world model in the contemporary Dreamer sense [Hafner et al., 2025, Ha and Schmidhuber, 2018]. The task is fully observed (state is the per-block 17-dim feature vector emitted by Tool 1; no raw observation), episodic (100-step cap, no partial observability), and the action space is finite categorical of size N_{blocks} (Methods, §9); we therefore do not learn a state representation from raw observations and do not use recurrent latent dynamics (RSSM). Combining a learned state encoder with the ranking-loss head reported here is a natural extension for settings where partial observability or pixel-level inputs apply (§9).

Per-parcel slope from DEM and cadastral polygons

The per-parcel slope feature used throughout the analysis is derived from the Third National Land Survey cadastral polygons and Copernicus DEM GLO-30 elevation tiles [European Space Agency, 2021], both freely available in their respective public archives. The DEM is kept in its native geographic CRS (EPSG:4326, 1-arcsecond resolution); slope is computed by central-difference gradients on the elevation grid using true east-west and north-south metres-per-pixel obtained from `pyproj.Geod` (WGS84 ellipsoid) at the study-area centre latitude. The pixel spacing is ~ 27 m east-west and ~ 31 m north-south at 29.5°N (the centre of both Bishan and Neijiang); reprojecting the DEM to a metric projection before gradient differencing would average over these non-square cells and systematically under-estimate slope. The slope raster, in degrees, is then aggregated to each cadastral polygon by area-weighted zonal mean over all DEM cells whose centre falls within the polygon. Polygons that contain no cell centre (typically very small parcels, $\lesssim 30$ m linear scale) inherit the per-county median slope to keep the downstream feature vector finite; on Bishan and Neijiang this affects $\sim 2\%$ and $\sim 3\%$ of parcels respectively, all in the smallest size class. The full prepare pipeline (DEM merge $\rightarrow$ slope $\rightarrow$ zonal mean $\rightarrow$ per-township block construction) takes ~ 3 minutes on Bishan and ~ 5 minutes on Neijiang on a 14-core desktop CPU and is released as the `farmland-mpc prepare` command (§Code Availability).

Problem setting and action-space encapsulation

We study county-level farmland consolidation in two real counties (Bishan District, Chongqing; Neijiang Dongxing District, Sichuan) and seven synthetic landscapes (§9). For Bishan, the planning area comprises 13 townships, 2,600 spatial blocks, and 52,515 cadastral parcels; for Neijiang, 29 townships, 3,711 blocks, and 76,376 parcels (Supplementary Table S1). Each block aggregates multiple parcels and is characterised by 17 features (farmland area, forest area, slope statistics, adjacency metrics, and investment indicators). Global context is encoded as an additional 12 features (remaining budget, contiguity, *baimu fang* count, etc.). The episode length is fixed at 100 steps for the real-data experiments and 50 steps for the synthetic benchmark (§5); each step has a budget of up to 5 paired farmland–forest swaps.

Agent action space. The agent’s action space is a single categorical choice over the spatial blocks per step: $a_t \in \{1, \dots, N_{\text{blocks}}\}$. The agent does *not* select individual parcels or the paired-swap sub-actions. Given the selected block, the environment internally executes *up to 5* paired farmland–forest swaps within that block using a deterministic rule-based procedure: candidate pairs are enumerated by joint feasibility under farmland–forest pairing and slope constraints, scored by a local improvement heuristic, and the top-scoring feasible pairs are applied until either 5 pairs are executed or the candidate pool is exhausted. The agent therefore chooses *where* to act, not *what* to do within the chosen location; this encapsulation keeps the action space at a tractable scale ($\sim 10^3$). The paired farm $\leftrightarrow$ forest swap conserves *counts* (one farmland parcel out, one forest parcel in) but not *areas*: parcels are paired by class and by within-block proximity, not by area, so any net cultivated-area change across an episode reflects the planner’s slope reward favouring high-slope (typically large) farmland exit and low-slope (typically small) forest entry. The unconstrained Bishan candidate scenario described in §4 loses 505 ha (1.05%) of total cultivated area as a consequence. This is not a hidden implementation detail; it is a failed hard constraint surfaced by the independent audit. The constrained runs reported in §4 activate `cultivated_area_floor_delta_ha`; with the default value 0.0, each committed pair must keep cumulative cultivated area at or above its initial value. A second execution floor, `baimu_area_floor_delta_ha`, can reject pairs that would

push cumulative qualifying *baimu fang* area below a user-specified threshold. The connectivity-conservative profile reported in Supplementary Table S17 uses the same action path but increases local pair-selection weights `gamma_conn` and `delta_conn` to favour swaps that preserve farmland adjacency during execution.

Action masking. Not all blocks are legal targets at every step. The action mask $m_t \in \{0, 1\}^{N_{\text{blocks}}}$ disables blocks that (i) have no remaining feasible parcel pairs under the local farm-forest and slope rules, (ii) are currently in a township whose per-township swap budget has been exhausted, or (iii) have been selected in the immediately previous step (a simple debounce to prevent action thrashing). MaskablePPO [Ross et al., 2011, Schulman et al., 2017, Huang and Ontañón, 2022]-style masking is applied in both the policy sampling and the MPC candidate enumeration, so that illegal blocks have zero probability or are excluded from the K candidate set. When the cultivated-area floor is active, the same mask additionally excludes any block for which no slope-improving within-block pair can be committed without violating the cumulative area floor; the greedy execution loop applies the same test again before every parcel-pair swap. The *baimu fang*-area floor is evaluated in the execution loop rather than in the full action mask because it requires a connected-component recomputation after tentative parcel-pair changes; if no candidate pair in a selected block can satisfy the floor, that block is removed from further execution in the current episode.

Reward formulation

Each step’s reward combines three normalised directional signals (slope, contiguity, and qualifying-patch area), a large-patch formation bonus, and an asymmetric area-protection penalty. Let $\bar{\sigma}_t$ be the area-weighted mean farmland slope at step t , C_t the within-farmland contiguity index, B_t the count of qualifying *baimu fang* (≥ 6.67 ha contiguous farmland patches), and A_t their total area. The per-step reward is

$$r_t = w_S \cdot \Delta \bar{\sigma}_t + w_C \cdot \Delta C_t + w_A \cdot \Delta A_t + b_B \cdot \max(0, B_t - B_{t-1}) + p_A \cdot \min(0, \Delta A_t) \quad (2)$$

where the directional changes are normalised by their initial county-wide values: $\Delta \bar{\sigma}_t = (\bar{\sigma}_t - \bar{\sigma}_0)/\bar{\sigma}_0$ (so that slope *reduction* contributes positively), $\Delta C_t = (C_t - C_{t-1})/C_0$, and $\Delta A_t = (A_t - A_{t-1})/A_0^{\text{farm}}$ (normalised by the initial total farmland area A_0^{farm} , so the weights act on relative rather than absolute changes). Weights $(w_S, w_C, w_A, b_B, p_A) = (4000, 500, 1500, 5, 2000)$ are inherited from our in-house DRL baseline configuration. The fourth term is a one-off bonus applied each time the qualifying-patch count increases between steps; the fifth term, $p_A \cdot \min(0, \Delta A_t)$, is an asymmetric penalty that activates only when qualifying area decreases, roughly doubling the marginal cost of shedding *baimu fang* area relative to the symmetric w_A reward (this asymmetry is what drives the agent to consolidate many small qualifying patches into fewer larger ones, increasing total qualifying *area* even as the qualifying *count* falls; §4). Under a median single-step decomposition, slope improvement contributes $\sim 60\%$, contiguity $\sim 15\%$, qualifying-patch area $\sim 20\%$, and the formation bonus $\sim 5\%$. We did not re-tune the weights for the present experiments; contrastive MPC, dream-trained PPO, and the model-free baselines all use the identical reward function, so any method-wise difference is attributable to the learning procedure rather than to reward engineering. A separate fixed penalty of -1 is applied whenever a selected block admits no feasible paired swap, discouraging wasted actions. Throughout the paper, “Slope %” is the relative change $\text{Slope \%} = (\bar{\sigma}_T - \bar{\sigma}_0)/\bar{\sigma}_0 \times 100$; contiguity and *baimu fang* count/area changes are absolute.

Block-level transition model

The transition model is a feed-forward neural network with approximately 237,000 parameters that takes as input the current block features (all blocks, 17 dimensions each), the global features (12 dimensions), and the selected action (block index), and predicts the next-step block features, global features, and scalar reward. Only the selected block changes (residual formulation). The base loss is

$$\mathcal{L}_{\text{MSE}} = \text{MSE}(\Delta \hat{b}, \Delta b) + \text{MSE}(\Delta \hat{g}, \Delta g) + 0.1 \cdot \text{MSE}(\hat{r}, r), \quad (3)$$

trained with Adam (learning rate 10^{-3} , weight decay 10^{-5}) on a 90/10 train/validation split with patience-6 early stopping.

Ensemble uncertainty estimation

We train 3–5 independently initialised transition models on the same dataset with different random seeds [Lakshminarayanan et al., 2017]. At inference, the ensemble produces mean predictions for next state and reward, plus the reward standard deviation σ_r across members. The dream environment penalises high-disagreement actions: $r_{\text{out}} = \bar{r} - \lambda_{\text{disag}} \cdot \sigma_r$ with $\lambda_{\text{disag}} = 0.5$.

Policy architecture: parcel-scoring policy

The policy is a dimension-invariant scoring network compatible with MaskablePPO. For each block i , the scorer concatenates the block’s 17-dimensional local features with the 12-dimensional global state, passes the result through a shared MLP ($[64, 32] \rightarrow 1$), and produces a scalar score. The action is the highest-scoring valid block. A separate value network takes only the global features and predicts the state value.

Dagger-style iterative refinement

We adopt a DAgger-style loop [Ross et al., 2011]: (1) collect an initial dataset D_0 of 6,000 transitions from random and greedy policies on the real environment; (2) train the transition model ensemble on $D_{\leq i}$; (3) train a policy in the dream environment for 50,000 timesteps; (4) evaluate on the real environment; (5) collect 10 stochastic on-policy episodes ($\sim 1,000$ transitions) and augment the dataset. Three iterations are performed.

Targeted evidence collection and stabilisation

In targeted DAgger, step (5) above is modified: each transition is scored by ensemble reward disagreement σ_r , and only transitions exceeding threshold $\tau = 0.5$ are retained; under this threshold, $\sim 99.6\%$ of transitions are filtered out, accounting for the dream-policy slowdown documented in §3. Three stabilisation mechanisms are applied across all protected runs. *Best-checkpoint selection*: every 10,000 PPO timesteps, the current policy is evaluated on the real environment (3 episodes); the checkpoint with the highest mean reward is retained. *Iteration-1 rollback*: if iteration-1 reward drops more than 50 points below iteration-0, the dataset and policy are reverted. *Early stopping*: training halts if validation loss does not improve for 6 consecutive epochs.

Model-predictive planning protocol

We use random shooting model-predictive control [García et al., 1989, Williams et al., 2017, Rubinstein, 1999]. At each real-environment step t with state s_t :

1. Sample K random action sequences $\{a_{t:t+H-1}^{(k)}\}_{k=1}^K$, each of length H , from the action-masked support.
2. For each sequence, roll forward through the ensemble mean model: $\hat{s}_{t+1}^{(k)}, \hat{r}_t^{(k)} = f_\theta(s_t, a_t^{(k)})$, continuing autoregressively for H steps.
3. Score each sequence by cumulative predicted reward: $R^{(k)} = \sum_{h=0}^{H-1} \hat{r}_{t+h}^{(k)}$.
4. Execute $a_t^{(k^*)}$ where $k^* = \arg \max_k R^{(k)}$ on the real environment.
5. Advance to s_{t+1} (real) and repeat.

For the synthetic-benchmark sweep (§5), the same protocol is applied with $H=5$, $K=50$, and greedy continuation (each subsequent step within a sequence picks the predicted-reward argmax over 50 candidates rather than continuing the random rollout). This is the configuration that becomes optimal under the contrastive ensemble (Table 1).

Aggregate fidelity vs within-state action ranking under MSE training

The decoupling described in §3 is a property of the loss. Let $\hat{r}_\theta(s, a)$ be the reward head of a learned model trained on (s, a, r) pairs from a distribution $\mu(s, a)$, and let

$$\mathcal{L}_{\text{MSE}}(\theta) = \mathbb{E}_{(s,a) \sim \mu} [(\hat{r}_\theta(s, a) - r(s, a))^2] \quad (4)$$

be the standard objective. Two distinct quantities can be evaluated on this learned model: *aggregate fidelity*, defined as the population MSE itself, and *within-state ranking accuracy*

$$\mathcal{A}_{\text{rank}}(\theta) = \mathbb{E}_{s \sim \mu_s} \mathbb{E}_{a_i, a_j \sim \mu(\cdot|s)} [\mathbf{1}[\text{sign}(\hat{r}_i - \hat{r}_j) = \text{sign}(r_i - r_j)]] , \quad (5)$$

which is the quantity model-predictive top- K selection actually consumes when scoring action menus at a single environment state.

The two quantities are not constrained against each other under MSE training. Consider a setting in which, at every state s , the action-conditioned reward distribution is symmetric around a state-dependent mean $\bar{r}(s)$ with state-marginal variance σ_s^2 and within-state action variance $\sigma_a^2(s)$. The learned predictor $\hat{r}_\theta^*(s, a) = \bar{r}(s)$, which assigns the conditional mean to every action regardless of its identity, achieves population MSE equal to $\mathbb{E}_s[\sigma_a^2(s)]$ (the irreducible state-conditional variance) and within-state ranking accuracy equal to 1/2 on every pair with non-zero true reward difference, identical to random guessing. A learned model can therefore minimise MSE within a constant of the Bayes optimum while remaining indistinguishable from random on the very ranking decisions that drive planning. Our 51.6% rank accuracy at cosine similarity 0.998 is empirically close to this regime: the predicted reward standard deviation across the action menu (Table 1) is compressed by an order of magnitude relative to the true within-state variance, exactly as the conditional-mean predictor would behave. This is a counterexample showing MSE *does not guarantee* ranking fidelity in regimes where $\sigma_a^2(s) \ll \sigma_s^2$, not a statement of statistical independence: in regimes where $\sigma_a^2(s)$ is large relative to σ_s^2 , MSE training will in practice acquire much of the ranking signal as a by-product (§6 demonstrates this empirically).

Contrastive training protocol

Pairwise data generation. We sample states by rolling out random trajectories in the real CountyLevelEnv. At each step, we (i) snapshot the full environment state (land-use assignments, swap counts, block availability caches, running slope/contiguity sums; 19 attributes total), (ii) sample 50 valid actions uniformly from the action mask, (iii) execute each action from the snapshotted state, recording the immediate reward, and (iv) restore the snapshot before advancing with a fresh random action. We continue across multiple episodes until 1,000 states are collected, yielding 50,000 ground-truth (s, a, r) evaluations. The pairwise dataset $\mathcal{D}_{\text{pw}}$ is saved as $(s_i, \{a_{i,j}\}_{j=1}^{50}, \{r_{i,j}\}_{j=1}^{50})$ for $i = 1, \dots, 1000$. Generation took 45 minutes on commodity CPU hardware. Snapshot/restore is a simulator-specific capability; replay-buffer substitutes are an open question (Limitations).

Loss function. The augmented training objective adds a margin ranking term:

$$\mathcal{L}_{\text{rank}} = \mathbb{E}_{(s_i, a_{i,j}, a_{i,k}) \sim \mathcal{D}_{\text{pw}}} [\max(0, -\sigma_{jk} \cdot (\hat{r}_{i,j} - \hat{r}_{i,k}) + m)]$$

where $\sigma_{jk} = \text{sign}(r_{i,j} - r_{i,k}) \in \{-1, 0, +1\}$ is the ground-truth pairwise label, $m = 0.1$ is the margin, and pairs with $\sigma_{jk} = 0$ are excluded. At each training step, we sample 10 action pairs per state from 50–100 randomly selected pairwise-dataset states. Total objective: $\mathcal{L} = \mathcal{L}_{\text{MSE}} + \lambda_{\text{rank}} \mathcal{L}_{\text{rank}}$.

Margin choice. The margin $m = 0.1$ is approximately $1/8$ of the empirical ground-truth reward standard deviation (0.811 on $\mathcal{D}_{\text{pw}}$). A coarse sweep of $m \in \{0.05, 0.1, 0.2\}$ on a held-out validation fold showed a flat optimum: ranking accuracy at $\lambda_{\text{rank}}=5.0$ varied within ± 1.5 percentage points across these three choices.

Training. We use 3-member ensembles trained with Adam (lr = 10^{-3} , weight decay 10^{-5}), batch size 256, for 15 epochs with patience 20. Each ensemble trains in ~ 25 minutes on commodity CPU hardware. We sweep $\lambda_{\text{rank}} \in \{0.0, 1.0, 5.0\}$ using an 80/20 train/validation split of the pairwise dataset. MPC evaluation uses $H=5$, $K=50$, greedy continuation with 50 candidates per continuation step.

Synthetic benchmark generator

The seven synthetic landscapes used in Section 5 are produced by a deterministic, seed-controlled generator that reads a YAML preset (terrain type, target block count, parcel-density distribution, land-use fractions, fragmentation length-scale, seed) and emits a directory of GeoJSON, CSV, GeoTIFF, and JSON files identical in schema to the real-data preprocessor’s outputs. The generator is fully deterministic given (preset, seed).

Eight-stage pipeline. (1) *Domain tessellation*: Voronoi tessellation of n_{blocks} seeds in a rectangle of suitable extent (size driven by parcels-per-block target). (2) *DEM synthesis*: 2D Perlin noise with multi-octave fractal terrain, preset-controlled amplitude and length-scale, clipped to non-negative. (3) *Slope derivation*: Horn’s algorithm on the synthetic DEM, sampled to each block as area-weighted mean and max. (4) *Parcel subdivision within blocks*: secondary Voronoi inside each block polygon using the parcels-per-block distribution. (5) *Land-use assignment*: a Gaussian random field thresholded into farmland / forest / other classes, with preset-controlled length-scale driving fragmentation. (6) *Feature computation*: 17-dim block feature vectors using the same code path as the real-data `county_env.py` (farmland area, forest area, slope statistics, adjacency features), guaranteeing environment compatibility. (7) *Adjacency graph*: planar adjacency from shared block

edges, with median degree close to the preset target. (8) *Posterior calibration check* (anchors only): compute initial slope, contiguity, *baimu fang* count, and *baimu fang* area, and verify within $\pm 50\%$ of the real anchor values; deviations are surfaced in a calibration delta report rather than silently accepted.

Dual-anchor calibration. Two of the seven presets, `bishan_clone` and `neijiang_clone`, are calibrated against Bishan and Neijiang respectively such that any single seed reproduces the four anchor initial-state metrics within $\pm 50\%$ of the real anchor. The other five presets vary morphology along terrain, size, and fragmentation axes (§5) but are not constrained against any real anchor; they are parameter-controlled stress tests rather than evidence of regional generalisation. Contiguity is the weakest calibration axis ($\sim 6\%$ on Bishan clone, $\sim 44\%$ on Neijiang clone), reflecting Voronoi-based geometries’ tendency to be more uniformly contiguous than real cadastre shaped by ridges, roads, and water channels.

Compatibility shim. A thin `synthetic_env_loader.py` reads the generator outputs (`blocks.csv` + `adjacency.json`) and returns the same data structure that `county_env.py` expects from real data; no environment-code changes are required. The synthetic and real evaluation regimes therefore use identical simulator semantics, ensuring that any cross-regime difference is attributable to the landscape, not to the engine.

Restoration-prioritisation env (cross-domain verification)

The cross-domain verification of §6 reuses the entire farmland pipeline (sample, train, plan) by swapping only the env class and the per-unit feature vector. The restoration env is a finite-action gymnasium environment over n_{units} planning cells; the action space is $\{0, 1, \dots, n_{\text{units}} - 1\}$, and “select unit a ” is a one-time monotonic assignment (no swap, no reversibility). The action mask excludes already-selected units and units whose individual restoration cost would exceed the remaining budget. The reward at step t when selecting unit a_t is a fixed linear combination

$$r_t = \sum_k w_k \cdot \phi_k(a_t) + w_{\text{conn}} \cdot \frac{|\{j \in \mathcal{N}(a_t) : j \text{ already selected}\}|}{|\mathcal{N}(a_t)|} - w_{\text{cost}} \cdot \frac{c(a_t)}{B}$$

where ϕ_k are per-unit attribute scores read from a CSV (risk index, water-quality priority, bio-diversity, etc.), w_k are weights specified in a per-case scenario configuration file, $\mathcal{N}(a_t)$ is the queen-contiguity neighbour set, $c(a_t)$ is the unit’s restoration-cost proxy, and B is the cumulative budget. State features $b \in \mathbb{R}^{n_{\text{units}} \times 17}$ are the standardised per-unit attributes (zero-padded to match $K_{\text{block}} = 17$ so the same network architecture as farmland applies); the global state $g \in \mathbb{R}^{12}$ encodes fraction-of-budget-remaining, fraction-of-units-selected, fraction-of-steps-elapsed, and per-component cumulative reward. Snapshot/restore is implemented by capturing the boolean selection mask, cumulative reward components, budget used, and step counter; the env is purely combinatorial and has no expensive-to-restore geometric state, so pairwise dataset generation runs in seconds rather than minutes.

For the Buchanan VA case ($n_{\text{units}} = 562$), per-unit features are derived from the OSMRE e-AMLIS national abandoned-mine-lands inventory [Office of Surface Mining Reclamation and Enforcement, 2019], USGS NHD flowline data [U.S. Geological Survey, 2024b], USGS 3DEP elevation [U.S. Geological Survey, 2024a], and Census TIGER county boundaries [U.S. Census Bureau, 2023], all freely available as public open data; episode length is 50 steps and the cumulative cost cap is set to \$200,000 in notional cost-proxy units. The cost-proxy field is derived as `restoration_cost_proxy`

= unit area in hectares (rescaled by 10), which yields a roughly e-AMLIS-comparable ranking among units but should not be interpreted as a real-world dollar reclamation cost; the e-AMLIS unfunded-cost field is sparse and unevenly populated across the inventory, and we therefore use a transparent area-based proxy and report results in proxy units rather than claim a calibrated cost-benefit analysis. For the synthetic case ($n_{\text{units}} = 420$), feature, adjacency, and reward generators are deterministic given a seed and are released alongside the synthetic benchmark of §5; episode length is 60 steps and the cost cap is 45,435 in synthetic units (no real-world correspondence). Both cases share the same MPC configuration as farmland ($H=5$, $K=50$, greedy continuation, scoring=reward) and the same training hyperparameters ($\lambda_{\text{rank}}=5.0$ for contrastive, margin 0.1, three-member ensembles per seed, five seeds for both the contrastive ($\lambda=5$) and the MSE-only ($\lambda=0$) conditions to ensure a fair head-to-head comparison.

The cumulative wall-clock for the full restoration verification (two cases, one prepare each, one sample each, six ensemble trains per case, six MPC plans, one random baseline per case) is approximately 30 minutes on a 14-core M-series Apple Silicon CPU running 3 ensembles in parallel; the cases are released as part of the open reproduction kit (§Code Availability).

Simulator-cost crossover protocol

The simulator-cost analysis of §7 treats the per-step wall cost of `env.step` as an independent variable. To isolate it from confounds (action space size, reward function, optimisation algorithm), we instrument the restoration env’s reward dispatcher with a configurable `time.sleep(delay_s)` call invoked exactly once per `env.step`. We then run two algorithms at matched 30s wall budgets across delays $d \in \{0, 1, 10, 50, 100\}$ ms.

Simulated-annealing baseline. We run a permutation-space SA over the candidate-unit action ordering. Each iteration: sample $i, j \in \{0, \dots, 79\}$ uniformly; swap positions i and j in the current ordering; re-roll out the full episode; accept the swap if it improves the cumulative reward, or with probability $\exp(\Delta r / T_t)$ otherwise; geometrically cool T_t at rate 0.995. Initial ordering is the priority-score-greedy ordering (this gives SA a strong warm start). The number of iterations actually completed within the 30s wall budget is the diagnostic: at delay 0 it reaches $\sim 9,000$, at delay 100 ms it reaches ~ 6 .

Learned-surrogate MPC baseline. We run `mpc.select_action` with random continuation, $H=5$, $K=50$, $\gamma=0.99$, scoring=reward, on each of the seeded RNG paths $\{1, 18, 35, \dots\}$ until the wall budget is exhausted. We report the best cumulative episode reward seen. The per-step ensemble forward pass is ~ 25 ms (one batched ONNX call evaluating all $K=50$ candidates in parallel) and is decoupled from the injected `env.step` delay; only the eventual environment commit step pays the simulator cost.

Farmland baseline measurement. The farmland Bishan datapoint (Figure 3, gold star) is measured on the actual `CountyLevelEnv` without any injected delay. The natural step cost is profiled as 17.9 ± 0.3 ms over a 5-step probe repeated three times (the variance is dominated by per-step *baimu fang* cache hits / misses; in the worst case where the union-find aggregation triggers, single-step cost rises to ~ 60 ms). Under a 180s wall budget rather than 30s (chosen because the wall cost per SA iteration on farmland is ~ 1.8 s, so 30s would only allow ~ 16 iterations and SA would not even reach its initial cooling regime), simulated annealing reaches reward 9.04 at $n_{\text{iter}} = 100$, random restart over 98 episodes reaches 30.85, and contrastive MPC reaches 86.74. The full driver script is released in `farmland_mpc/tests/farmland_baselines.py`.

ArcGIS Pro toolchain

The deployment artefact described in §8 is an ArcGIS Pro Python Toolbox (`LandUseOptimization_P9.pyt`) exposing five tools: (1) *Prepare Data & Blocks* ingests a Third National Land Survey DLTB shapefile and a digital elevation model and emits the block-level state expected by the learned-model code; (2) *Sample Transitions* generates the pairwise dataset $\mathcal{D}_{\text{pw}}$ on the user’s data via the snapshot/restore protocol described above; (3) *Train Contrastive Ensemble* trains a 3-member ensemble at $\lambda_{\text{rank}}=5.0$ and exports each model to ONNX with the user’s N_{blocks} baked in; (4) *MPC Planning* runs the planning loop and writes back an optimised cadastral shapefile readable by any standard GIS workflow; (5) *Check Dependencies* validates the user’s Python environment (ArcGIS Pro 3.6+, Spatial Analyst extension, `gymnasium`, `libpysal`) before any pipeline run is permitted. Tools 1–3 are run once per county on commodity desktop hardware ($\sim 1\text{--}2$ hours total wall time); Tool 4 is the iterated planning loop (one ~ 7 -minute scenario per call). A 70-second smoke test on a 2-township subset of Neijiang Dongxing District completes the full pipeline as a continuous-integration check.

Three operational caveats are visible in the toolchain. First, the ONNX export bakes in N_{blocks} , so a pre-trained ensemble cannot be reused across counties; Tool 4 carries an explicit pre-flight check that refuses to run if the pre-trained ensemble’s N_{blocks} does not match the prepared dataset. Second, the default projected coordinate reference system (EPSG:32648, UTM zone 48N) covers central and western China; users in eastern, north-eastern, or north-western China must override the projection in Tool 1 to avoid slope errors. Third, Tool 4’s default mode optimises the stated reward and emits an auditable candidate; users who require no net cultivated-area loss should activate the hard cumulative floor parameter `cultivated_area_floor_delta_ha=0`. Users can additionally enforce no net qualifying *baimu fang* area loss with `baimu_area_floor_delta_ha=0` or bias execution toward local connectivity with `gamma_conn` and `delta_conn`. The end-to-end CLI Bishan replication on the package-side prepare achieves a slope change of $-1.5439 \pm 0.0409\%$ (1 shipped ensemble, 5 episodes), bracketed by the lab-pipeline matched-DRL figure of $-1.289 \pm 0.079\%$ (5 independently trained ensembles, 1 episode each, §4) and the on-disk GIS verification anchor of -2.001% (§9); the pipeline-to-pipeline gap reflects the corrected geographic-CRS slope step in the package-side prepare. With the no-net-loss floor active, the canonical Bishan deployment run reaches -0.675% slope and $+3.03$ ha cultivated area in 300 s, while the corresponding Neijiang run reaches -0.492% slope and $+62.2$ ha cultivated area in 478 s. A complete cross-pipeline reconciliation table is reported in Supplementary Section S3, and the seven-profile Bishan execution-constraint frontier is reported in Supplementary Table S17.

Evaluation protocol

For the real-data baseline-comparison experiments (§4), each stochastic configuration is evaluated with 5 deterministic episodes or 5 independently trained ensembles on the real `CountyLevelEnv`, as stated next to each result. The GIS-only unconstrained audit, the hard no-net-loss constrained runs, and the Bishan execution-constraint frontier are single canonical deployment episodes, because their purpose is to verify on-disk Tool 4 outputs and policy status rather than estimate training-seed variance. Reported metrics include reward, slope change (%), contiguity change, *baimu fang* count change, total *baimu fang* area change, steep-tail area change, and total cultivated-area change where relevant. For the synthetic benchmark (§5), each method is run for 5 random seeds on each of the 7 presets; per-cell results are aggregated as cross-seed mean $\pm$ sample standard deviation. Cross-seed statistics for the contrastive configurations are computed across 5 independently trained ensembles (different random seeds for both the pairwise dataset and the network initialisation). Tests against

fixed reference values use one-sample t -tests; superiority comparisons between two methods on independent seeds use Welch’s two-sample t -tests; superiority comparisons between two configurations on matched seeds use paired t -tests, and equivalence at a stated reward-units margin uses two-one-sided-tests (TOST); ranking-distribution comparisons additionally use Wilcoxon signed-rank as a robustness check. Reported quantities are test type, df, p , and Cohen’s d (or d_z for paired tests).

Independent GIS-grounded verification of reported gains

Because the primary metrics in §4 are computed inside the same `CountyLevelEnv` whose learned ensemble drives the MPC scoring loop, a sceptical reviewer can ask whether the reported slope reduction is a real change in the optimised cadastre or an artefact of the learned dynamics agreeing with the planner. To resolve this, we recompute all primary metrics directly from the optimised DLTB shapefile written to disk by Tool 4, using a script that depends only on `geopandas` and `libpysal.weights.Queen` and shares no source with the optimisation pipeline (no `torch`, no `ONNX`, no `env` class). The script reads the per-parcel `CHG_FLAG` audit field that Tool 4 emits, reconstructs the post-optimisation land-use vector, joins per-parcel `slope_mean` from the prepared DEM-derived raster on `BSM`, reprojects to EPSG:32648 for area computation, builds Queen contiguity, and applies the metric definitions exactly as stated in §2 to both the original and the optimised land-use vectors, including total cultivated area. This verification operates on the package-side prepare (Tool 1 output, 2,640 blocks, 53,004 swappable parcels) and a single canonical episode of the shipped ensemble.

On the 53,086-parcel Bishan deployment dataset, the recomputed deltas (independent path) and the env-reported deltas (planning path) agree to floating-point reduction noise: slope -2.0000% vs. -2.0006% ($|\Delta|=6\times10^{-4}$ pp); contiguity $+0.012739$ vs. $+0.012748$ ($|\Delta|=8\times10^{-6}$); *baimu fang* area -473.30 ha vs. -473.30 ha ($|\Delta|<10^{-4}$ ha); and the *baimu fang* count change reproduces exactly at $+8$ patches. The 428 farm→forest and 428 forest→farm swaps written to the shapefile are also bit-identical to the env’s tally, confirming the farmland-count conservation invariant on disk. For the constrained Bishan run, the same independent script reports slope -0.6746% , contiguity $+0.02037$, *baimu fang* area $+29.80$ ha, and cultivated area $+3.03$ ha; the differences from the planner-reported values are 1.5×10^{-4} pp on slope, 6.4×10^{-5} on contiguity, and 0 ha on *baimu fang* area. The reported slope, contiguity, and cultivated-area changes are therefore properties of the on-disk geometry, not of the learned dynamics model.

We disclose two transparency caveats. First, the lab-pipeline five-ensemble sweep cited in §4 is by construction deterministic at planning time: under fixed ensemble member ordering, top- K rollout scoring, and greedy continuation, the seed initialises RNGs that are not reached on the action-selection path; the cross-seed standard deviation in that comparison comes from the five *independently trained ensembles*, not from re-rolling a single ensemble. To surface a meaningful variance estimator on the package-side single-ensemble protocol used here, we replaced seed jitter with *ensemble member subset ablation*, re-running MPC three times under the three distinct 2-of-3 member subsets. Across these replicates the slope reduction is $-1.998 \pm 0.026\%$ (range $[-2.022, -1.971]$), contiguity change is $+0.0131 \pm 0.0016$, and *baimu fang* area change is -474.5 ± 31.1 ha—sandwich-bracketing the full-ensemble result of -2.001% , $+0.01275$, -473.3 ha to within sample standard deviation on every metric. The full ensemble therefore sits at the median of subset performance rather than benefiting from any single member, and operational variability under the strongest available counter-factual (a different 2-of-3 subset) is bounded at ± 0.026 pp on slope. Second, the verification script keeps 53,086 swappable parcels (those with both a slope value and a farm/forest `ORIG_DLEB`), 82 more than the env’s 53,004; the excess corresponds to parcels whose township code falls outside the configured registry, all of which carry `CHG_FLAG` = 0 and therefore do not affect

the delta computation. The verification script is released alongside the toolchain (see *Code availability*) so that any third party with Tool 4’s optimised shapefile and a slope raster can reproduce the reported numbers without invoking the optimisation pipeline.

Direct measurement of ensemble 1-step prediction accuracy

The on-disk verification above only rules out errors in the writer or the metric definitions; it does not bound the contribution of dynamics-model error to the planner’s choices. To address that, we rolled the same MPC policy ($H=5$, $K=50$, $\gamma=0.99$, greedy continuation, reward scoring) for 100 steps on the Bishan deployment dataset and recorded, at every step, both the ensemble’s prediction and the true `env.step` output. The action sequence and per-step metrics reproduce episode 0 of §4 to floating-point identity (slope at steps 10/20/.../100 matches to four decimals), confirming determinism of the deployed pipeline.

Across the 100 in-distribution steps, the ensemble’s reward predictions have MAE 0.818 against an empirical $\sigma(r_{\text{true}}) = 1.40$, with predictions confined to a narrow band of width 0.24 (range $[+0.43, +0.67]$) while the true reward varies in $[-3.04, +6.36]$. The ensemble is therefore strongly mis-calibrated as a regressor, scoring barely better than a constant predictor (naive-mean baseline MAE 0.83). Despite this, the Spearman rank correlation between predicted and true reward is +0.22, and this is what MPC’s top- K selection consumes; the trajectory match to §4 confirms that ranking signal is sufficient at $H=5$ horizon. Next-state predictions for the slope-improvement channel (`gf[4]`, used directly when `scoring="slope"`) carry MAE 2.4×10^{-3} , smaller than the channel’s own per-step standard deviation (5.2×10^{-3}). These three numbers together are a direct empirical instantiation of the ranking-versus-regression decomposition we developed in §3: the contrastive ensemble has not learned to predict reward magnitude, but it has learned to rank candidate actions, and the planner’s $H=5$ top- K mechanism converts that weak ranking signal into a deterministic, physically-correct trajectory.

A third transparency caveat follows from this measurement: ensemble disagreement is not informative as an uncertainty signal under the current bootstrap diversity. The mean of r_{std} is 0.027 across the 100 steps, the Pearson correlation between r_{std} and the per-step absolute reward error is +0.08, and the maximum disagreement seen (0.075) is two orders of magnitude smaller than the maximum absolute error (5.07). Members under-fit the reward distribution similarly to one another, so disagreement does not flag the cases where they jointly are wrong. Practitioners who want disagreement-driven candidate filtering should train Tool 3 with stronger diversity (different feature subsets or training data slices, not just initialisation seed). The default `scoring="reward"` pathway in Tool 4 remains safe because top- K ranking, not absolute reward, drives planning; `scoring="slope"` is an alternative that bypasses the reward head entirely and uses the better-calibrated `gf[4]` channel.

Replacing the ensemble with the true environment

As a final ablation we re-ran the same MPC procedure ($H=5$, $K=50$, $\gamma=0.99$, single rollout per candidate, seed 0) but used the actual `env.step` in place of every ensemble call. Two implementation concessions are required to make this tractable: stage 1 ranks a random subset of 200 valid actions per step instead of all $\sim 2,000$ (the ensemble’s `batch_predict` scores every action at near-constant marginal cost; a true-env stage 1 must run incremental simulation per action), and stage 2 uses random continuation instead of greedy. Mutable env state (`land_use`, `swapped`, `farmland_nbr_count`, accumulators, baimu cache) is snapshotted before each trial step and restored after; the *baimu fang* union-find is frozen inside scoring trials with the outer commit step computing it correctly. Total

wall time was 358 s on the same hardware that ran the 1,448 s ensemble baseline.

The resulting trajectory finishes at slope -1.819% , contiguity $+0.0087$, and *baimu fang* area change -565 ha, all three strictly worse than the package-side ensemble pipeline’s -2.001% , $+0.0128$, -473 ha. Per-step inspection shows true-env MPC starting behind (the 200-action stage 1 misses the strongest moves the ensemble identifies via full-action coverage), narrowing the gap as the trajectory matures, and finishing 0.18 pp short of the ensemble’s slope. The reason is structural rather than statistical: ensemble `batch_predict` grants the planner the depth and breadth that incremental simulation cannot afford within the same compute budget. Putting matched effort into true-env MPC—full enumeration in stage 1 and greedy continuation in stage 2—would cost approximately five days of wall time per episode for a payoff of at most a fraction of a percentage point of slope. The reported -2.00% candidate is therefore on the optimistic side of the planner’s accessible frontier under matched compute, and the ensemble’s contribution to the deployment is structural (it grants the planner depth and breadth) rather than illusory.

Data availability

The raw farmland cadastral records for Bishan District (Chongqing) and Neijiang Dongxing District (Sichuan) are derived from the Third National Land Survey of China and cannot be publicly redistributed under existing data-governance restrictions. Researchers requiring access to the raw cadastral records should approach the Chongqing and Sichuan provincial natural-resources departments under the same procedure used for any access to Third Land Survey data; we have no role in granting access. To support reproducibility and verification independent of restricted access, all derived data products required to reproduce the reported experiments are publicly available now (not gated on publication) at the project repository (§Code Availability) under MIT licence:

(a) the full synthetic farmland benchmark of seven landscape presets with deterministic generators (Section 5) under CC-BY 4.0; this benchmark alone is end-to-end sufficient to reproduce the contrastive-MPC pipeline’s qualitative findings and to benchmark future methods without any access to restricted data; (b) aggregated block-level features (the 17-dimensional vectors for all 2,600 Bishan and 3,711 Neijiang blocks in the initial state) and the anonymised pairwise dataset $\mathcal{D}_{pw}$ for both farmland counties, sufficient to reproduce transition-model training, contrastive training, and MPC evaluation end-to-end without raw parcel records; (c) the released contrastive ensembles for both Bishan and Neijiang (5×3 trained members per case, plus the bishan λ -ablation set used in §3); (d) random seeds, hyperparameters, and de-identified training logs for all reported farmland experiments; (e) the complete restoration cases of §6: *Buchanan County, Virginia, USA* (planning-unit attributes, adjacency, OSMRE e-AMLIS extracts, USGS NHD flowlines clipped to the county, USGS 3DEP slope, and Census TIGER county boundary, all already public per their respective US-government open-data licences) and the deterministic *synthetic* mine-restoration generator with all seven weight-sweep Pareto configurations.

The synthetic farmland benchmark (a) and the public-data Buchanan restoration case (e) jointly form a fully open reproduction track that requires no restricted-data access; together with the released ensembles (c), an external reviewer or future user can reproduce all qualitative findings of this paper end-to-end on commodity desktop hardware. The cadastral derivatives (b)–(d) extend that reproduction to the actual real-county Bishan and Neijiang results without requiring the raw cadastral parcels themselves, under the same data-governance constraints that apply to many Earth science datasets covering politically sensitive regions.

Code availability

All code for the contrastive training, MPC planner, synthetic generators, ArcGIS Pro toolchain, the non-commercial CLI workflow, the operations-research baselines (§6 Table 2: greedy, simulated annealing, NSGA-II, CBC-MILP), the planner-relevant ranking metrics (NDCG, top- K regret), the restoration Pareto-sweep driver, and the Bishan execution-constraint frontier driver is released as a public repository on GitHub at the time of submission (<https://github.com/zhouning/arcgis-farmland-mpc>), under an MIT licence, with step-by-step reproduction instructions in `docs/REPRODUCE.md` and a smoke test that completes a full Tool 1–4 cycle in ~ 70 seconds on a small subset.

Author contributions

N.Z. conceived the study, designed and implemented the methods, conducted all experiments, and wrote the manuscript. X.J. supervised the research and revised the manuscript.

Competing interests

The authors declare no competing interests.

Acknowledgements

This work was completed during N.Z.’s graduate study at the School of Software and Microelectronics, Peking University, and the deployment toolchain (Section 8) was developed at the Public R&D Center, SuperMap Software Co., Ltd.

Funding

This research received no external funding.

ORCID

Ning Zhou 0009-0002-5647-7388

References

- Jeroen C. J. H. Aerts and Gerard B. M. Heuvelink. Using simulated annealing for resource allocation. *International Journal of Geographical Information Science*, 16(6):571–587, 2002.
- Christopher Burges, Tal Shaked, Erin Renshaw, Ari Lazier, Matt Deeds, Nicole Hamilton, and Greg Hullender. Learning to rank using gradient descent. In *Proceedings of the 22nd International Conference on Machine Learning (ICML)*, pages 89–96, 2005. doi: 10.1145/1102351.1102363. URL <https://www.microsoft.com/en-us/research/publication/learning-to-rank-using-gradient-descent/>. Seminal pairwise learning-to-rank work (RankNet; cross-entropy over pairwise preferences).

- Christopher J. C. Burges. From RankNet to LambdaRank to LambdaMART: An overview. Technical report, Microsoft Research Technical Report MSR-TR-2010-82, 2010.
- Kai Cao, Bo Huang, Shaowen Wang, and Hui Lin. A boundary-based generic genetic algorithm for sustainable land-use planning. *Computers, Environment and Urban Systems*, 36(3):257–269, 2012. doi: 10.1016/j.compenvurbsys.2011.08.001. CEUS, genetic-algorithm baseline for sustainable land-use planning; representative pre-DRL method on closely related task.
- Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems*, volume 30, 2017. URL <https://papers.nips.cc/paper/2017/hash/d5e2c0adad503c91f91df240d0cd4e49-Abstract.html>.
- Kurtland Chua, Roberto Calandra, Rowan McAllister, and Sergey Levine. Deep reinforcement learning in a handful of trials using probabilistic dynamics models. In *Advances in Neural Information Processing Systems*, volume 31, 2018. URL <https://papers.nips.cc/paper/2018/hash/3de568f8597b94bda53149c7d7f5958c-Abstract.html>.
- Kalyanmoy Deb, Amrit Pratap, Sameer Agarwal, and T. A. M. T. Meyarivan. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation*, 6(2):182–197, 2002.
- Julian Equihua et al. Connectivity conservation planning through deep reinforcement learning. *Methods in Ecology and Evolution*, 2024. DRL applied to connectivity-aware conservation; related precedent for combinatorial spatial planning with ecological objectives.
- European Space Agency. Copernicus DEM – global and European digital elevation model (COP-DEM), 2021. Product handbook GLO-30, version 1.1; available via <https://copernicus-dem-30m.s3.amazonaws.com/>.
- Jonathan A. Foley, Navin Ramankutty, Kate A. Brauman, Emily S. Cassidy, James S. Gerber, et al. Solutions for a cultivated planet. *Nature*, 478(7369):337–342, 2011.
- Carlos E. García, David M. Prett, and Manfred Morari. Model predictive control: Theory and practice—a survey. *Automatica*, 25(3):335–348, 1989.
- David Ha and Jürgen Schmidhuber. Recurrent world models facilitate policy evolution. In *Advances in Neural Information Processing Systems*, volume 31, 2018. URL <https://papers.nips.cc/paper/2018/hash/2de5d16682c3c35007e4e92982f1a2ba-Abstract.html>.
- Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse domains through world models. *Nature*, 640:647–653, 2025. doi: 10.1038/s41586-025-08744-2. URL <https://www.nature.com/articles/s41586-025-08744-2>.
- Nicklas Hansen, Hao Su, and Xiaolong Wang. TD-MPC2: Scalable, robust world models for continuous control. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://openreview.net/forum?id=0xh5CstDJU>. Spotlight.
- Ralf Herbrich, Thore Graepel, and Klaus Obermayer. Large margin rank boundaries for ordinal regression. In *Advances in Large Margin Classifiers*, pages 115–132. MIT Press, 2000. Original large-margin pairwise ranking formulation for ordinal regression.

- Jikun Huang, Wei Wei, Qi Cui, and Wei Xie. The prospects for china’s food security and imports. *China Agricultural Economic Review*, 2018.
- Shengyi Huang and Santiago Ontañón. A closer look at invalid action masking in policy gradient algorithms. In *The International FLAIRS Conference Proceedings*, 2022.
- Michael Janner, Justin Fu, Marvin Zhang, and Sergey Levine. When to trust your model: Model-based policy optimization. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://papers.nips.cc/paper/2019/hash/5faf461eff3099671ad63c6f3f094f7f-Abstract.html>.
- Krzysztof Janowicz, Song Gao, Grant McKenzie, Yingjie Hu, and Budhendra Bhaduri. GeoAI: Spatially explicit artificial intelligence techniques for geographic knowledge discovery and beyond. *International Journal of Geographical Information Science*, 34(4):625–636, 2020. doi: 10.1080/13658816.2019.1684500. GeoAI overview; situates AI methods in geographic knowledge discovery.
- Thorsten Joachims. Optimizing search engines using clickthrough data. In *Proceedings of the 8th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 133–142, 2002. doi: 10.1145/775047.775067. Ranking SVM; large-margin pairwise ranking in its hinge formulation—the direct ancestor of our margin ranking loss.
- Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems*, volume 30, 2017. URL <https://papers.nips.cc/paper/2017/hash/9ef2ed4b7fd2c810847ffa5fa85bce38-Abstract.html>.
- Y. Liu and Y. Li. Farmland fragmentation and its impact on agricultural productivity: Evidence and implications for land consolidation policy. *Land Use Policy*, 2023. Domain motivation for land consolidation; fragmentation–productivity link.
- Yansui Liu, Fang Fang, and Yuheng Li. Land consolidation projects in rural china: Origins, progress, and outcomes. *Land Use Policy*, 40:3–14, 2014.
- Hualou Long. Land use policy in china: Introduction. *Land Use Policy*, 40:1–2, 2014.
- S. Niu et al. Quantifying farmland fragmentation: Indicator system and spatial–temporal patterns. *Land*, 2023. [v7-PLACEHOLDER] Verify full author list, volume, issue, and DOI before submission. Used in v7 Section 1 to support “per-hectare labour cost, mechanisation barriers, climate-driven yield-shock vulnerability” framing. Aligned with the FFI implementation in the author’s own FFI.py.
- Office of Surface Mining Reclamation and Enforcement. Abandoned mine land inventory system (e-AMLIS), 2019. Public ArcGIS REST service; inventory snapshot September 2019; <https://amlis.osmre.gov/>.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744, 2022. URL <https://papers.nips.cc/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html>.

- Markus Reichstein, Gustau Camps-Valls, Bjorn Stevens, Martin Jung, Joachim Denzler, Nuno Carvalhais, and Prabhat. Deep learning and process understanding for data-driven Earth system science. *Nature*, 566(7743):195–204, 2019.
- David Rolnick, Priya L. Donti, Lynn H. Kaack, Kelly Kochanski, Alexandre Lacoste, et al. Tackling climate change with machine learning. *ACM Computing Surveys*, 55(2):1–96, 2022.
- Rosalaura Romeo, Sara Manuelli, Mariella Geringer, and Valeria Barchiesi. Mountain agriculture: opportunities for harnessing zero hunger in Asia. FAO report, 2020. [v7-PLACEHOLDER] Verify exact title, authorship, year, and pagination before submission. Used in v7 Section 1 and Abstract framing for “roughly a fifth of global agricultural production from mountainous landscapes”. If the precise figure cannot be sourced from this report, drop the quantification or replace with a Mountain Partnership statistical reference.
- Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 15, pages 627–635. PMLR, 2011. URL <https://proceedings.mlr.press/v15/ross11a.html>.
- Mark D. A. Rounsevell, Bas Pedroli, Karl-Heinz Erb, Marc Gramberger, Anne Gravsholt Busck, et al. Challenges for land system science. *Land Use Policy*, 29(4):899–910, 2012.
- Reuven Y. Rubinstein. The cross-entropy method for combinatorial and continuous optimization. *Methodology and Computing in Applied Probability*, 1(2):127–190, 1999.
- Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Mastering Atari, Go, chess and shogi by planning with a learned model. *Nature*, 588(7839):604–609, 2020. doi: 10.1038/s41586-020-03051-4. URL <https://www.nature.com/articles/s41586-020-03051-4>.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, et al. Mastering the game of go without human knowledge. *Nature*, 550(7676):354–359, 2017.
- Marco Springmann, Michael Clark, Daniel Mason-D’Croz, Keith Wiebe, Benjamin Leon Bodirsky, et al. Options for keeping the food system within environmental limits. *Nature*, 562(7728): 519–525, 2018.
- Theodor J. Stewart, Ron Janssen, and Marjan van Herwijnen. A multiobjective genetic programming approach to land use planning. *Computers & Operations Research*, 31(14):2293–2313, 2004.
- Devis Tuia, Konrad Schindler, Begum Demir, Xiao Xiang Zhu, Mrinalini Kochupillai, et al. Artificial intelligence to advance Earth observation: A review of models, recent trends, and pathways forward. *IEEE Geoscience and Remote Sensing Magazine*, 2024.
- U.S. Census Bureau. TIGER/line cartographic boundary files, 2023. 500K-resolution county boundaries; <https://www.census.gov/geographies/mapping-files/>.
- U.S. Geological Survey. 3d elevation program (3DEP), 2024a. Elevation imagery and derived slope rasters; <https://www.usgs.gov/3d-elevation-program>.

- U.S. Geological Survey. National hydrography dataset (NHD), 2024b. Hydrography MapServer; <https://hydro.nationalmap.gov/arcgis/rest/services/nhd/>.
- Grady Williams, Paul Drews, Brian Goldfain, James M. Rehg, and Evangelos A. Theodorou. Model predictive path integral control: From theory to parallel computation. *Journal of Guidance, Control, and Dynamics*, 40(2):344–357, 2017.
- Y. Zheng, Y. Lin, L. Zhao, et al. Reinforcement learning for community spatial planning. *Nature Computational Science*, 2023. Deep-RL applied to urban community spatial planning; reference for DRL in structured discrete spatial design.